

Chapter 1

Prelude

Chapter 2

Introduction

blah blah blah

2.1 Sequential versus parallel computing

Developers worldwide are, and have always been, on the lookout for increased computing performance.

Until recently, the increased performance could easily be achieved through advances within raw computing power, as CPU's had steadily been doubling their number of on-chip transistors, in rough accordance to Moore's Law (citér her). However, the performance increases in CPU design has now slowed down significantly, due to physical limitations to CPU design.

Instead of adding more transistors or increasing the clock frequency of newer CPUs, the CPU manufacturers have instead opted to split their single-core CPUs up into multicore CPUs, which means that any program can now run several threads on the CPU's cores, simultaneously. The CPU's cores are specialized in advanced computations 's cores are

In sectors like the financial sector and within the natural sciences, there is a need for handling large amounts of data in an effective manner. With algorithmic trading gaining ground within the trading sector, the trader who can analyze incoming buy- and sell-offers the fastest, usually has the advantage at the exchanges. Likewise, faster computing can increase productivity for chemists and biologists who are analysing large datasets, physicists can run faster simulations, and so on. All of these activities are based on executing relatively simple computations on enormous datasets. The hundreds of simultaneous threads on the GPUs are, compared to the CPU, optimal for performing these calculations as fast as possible, which is why GPUs are increasingly being used for *General Purpose Computing on Graphics Processing Units*.

2.2 Parallel programming in Futhark

GPU programming is in principle easily available for everyone. As long as the user has access to a GPU and a reasonable PC for developing software, it just takes a bit of effort and reading to get started with CUDA, OpenCL or similar

programming. Realistically however, it takes much more than just a little effort to start writing one's own GPU programs.

Take the function $f(x) = ax + y$. In figure 2.1 we see the function implemented as a CUDA program. In this program, we are defining the kernel `saxpy` itself, and also manually copying data back and forth between the GPU. Now take the same program, written in Futhark as in figure 2.2.

Whereas the CUDA kernel needs to check whether the current thread is outside of the bounds of the input data, the equivalent kernel in Futhark is simply a declaration of its function. Also, the Futhark version does not bother with getting array elements by the current thread ID.

In the main function itself, the initial lists are generated by functions, and the user doesn't have to allocate space on neither the computer *host*, or the GPU *device*. The hard work is done by a Futhark SOAC, which is eventually compiled into a kernel

Of course, Futhark is compiled into either C- or Python code that does indeed contain `malloc` calls, GPU kernels with bounds checking, and so on, but that part is very well hidden from the Futhark programmers themselves. All in all, writing effective GPU programs becomes much more accessible when it's possible to do in a declarative manner, like Futhark, without also having to the minute details that comes with GPU computations.

2.3 Motivation

FShark is intended to be a way of writing and utilizing Futhark, without actually having to write or interact with the Futhark language and compiler itself. Besides some tooling and an F# SOAC library, it primarily consists of the FShark compiler that compiles from F# source code to Futhark source code, and the Futhark C# generator, which compiles Futhark programs as either standalone C# programs or -libraries.

As much as most developers are happy to increase performance on big computations, it is not always an option to incorporate an extra language into an already existing programming language. At this moment, using Futhark in either a C#- or F# project is a contrived process that usually requires spawning a subprocess with a `futhark-opencl` C program from inside one of the .NET projects.

In order to use Futhark natively in .NET languages, it is therefore necessary to write a backend for Futhark in a .NET language. For FShark, I have chosen to implement this backend in C#, as the Futhark intermediate code `ImpCode`¹ is trivial to translate into imperative C# statements and expressions. Also, there are C# libraries available which supply OpenCL bindings, which are needed to implement the necessary OpenCL constructs from `ImpCode`.

It is my belief that exporting Futhark programs as .NET executables and -libraries will lower the barrier to Futhark usage in .NET projects significantly, hopefully increasing the all-round number of Futhark users, and in the long term, increasing utilization of GPU programming and making it more widely available.

However, one could do even more than just exporting Futhark to .NET, to increase accessibility:

¹which stands for Imperative Code

```

#include <stdio.h>

__global__
void saxpy(int n, float a, float *x, float *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}

int main(void)
{
    int N = 1<<20;
    float *x, *y, *d_x, *d_y;
    x = (float*)malloc(N*sizeof(float));
    y = (float*)malloc(N*sizeof(float));

    cudaMalloc(&d_x, N*sizeof(float));
    cudaMalloc(&d_y, N*sizeof(float));

    for (int i = 0; i < N; i++) {
        x[i] = 1.0f;
        y[i] = 2.0f;
    }

    cudaMemcpy(d_x, x, N*sizeof(float), cudaMemcpyHostToDevice);
    cudaMemcpy(d_y, y, N*sizeof(float), cudaMemcpyHostToDevice);

    // Perform SAXPY on 1M elements
    saxpy<<<(N+255)/256, 256>>>(N, 2.0f, d_x, d_y);

    cudaMemcpy(y, d_y, N*sizeof(float), cudaMemcpyDeviceToHost);

    cudaFree(d_x);
    cudaFree(d_y);
    free(x);
    free(y);
}

```

Figure 2.1: $ax + y$ in CUDA

```

let saxpy (a : f32) (x : f32) (y : f32) : f32 =
    a*x+y

let main =
    let N = 1<<20
    let a = 2f32
    let xs = replicate N 1f32
    let ys = replicate N 2f32
    let ys = map2 (saxpy a) xs ys
    in ys

```

Figure 2.2: $ax + y$ in Futhark

As tens of thousands of programmers worldwide (CHECK NUMBER JEEEEZ) are already writing F# programs, and that most of F#'s functional language features can be directly translated into equivalent Futhark features, it became worthwhile to investigate whether it was possible to design a way for users to both write and utilize Futhark in F# projects, without ever actually touching the Futhark language or compiler themselves. Instead, users can write their data parallel F# modules in FShark, and compile these modules into Futhark libraries automatically.

In this case, it would be possible to get Futhark speeds in F# programs, without doing much more than installing the Futhark compiler locally, and adding the required FShark libraries to the F# project.

It is my belief that being able to achieve Futhark performance in regular F# programs almost automatically, will make it significantly easier for people to adapt to Futhark programming.

(SOME MORE MORE SOME MORE)

2.4 The contributions of this thesis

The contributions of this thesis are as follows:

1. The FSharkPrelude:

The FSharkPrelude is a subset of the Futhark SOACs, ported to F#. To write an FShark program, the user is directed to limit himself to the SOACs in the FSharkPrelude. This means exchanging `Array.map` for `FSharkPrelude.Map`, `Array.foldBack` for `FSharkPrelude.foldr`, and so on. However, the FSharkPrelude carries the guarantee, that all FSharkPrelude functions works equivalently to their Futhark SOAC namesakes. This prelude, together with the F# subset chosen for FShark, makes it possible to write F# programs which, when translated to Futhark code, are equivalent to their Futhark counterparts.

2. An F# subset translatable with FShark:

As F# is not only a multi-paradigm language, but also has access to the entire standard .NET library, it was required to make FShark support only

a subset of F#. This has been implemented by whitelisting only the F#-to-Futhark translatable types, constructs and expressions in the FShark compiler. Furthermore, no other module imports than FSharkPrelude are allowed. This subset is of course documented for users.

3. The FShark Compiler and Wrapper:

The FShark Compiler and Wrapper takes a module written in FShark as input, converts the FShark module into a compiled Futhark C# module, and makes it available to the F# program, all at runtime. The pipeline is described in sec ??

4. A C# backend for Futhark:

To actually use Futhark in C# projects (and transitively F# projects), it was necessary to develop and add a C# backend to the Futhark compiler. This backend is equivalent in functionality to the C- and the Python backends that are already available.

2.5 Roadmap

The main part of this thesis is split in four parts. blaaaah

Chapter 3

Background

FShark is built on the interaction between Futhark, F# and C#, wherefore WE SHOULD HAVE A BETTER INTRODUCTION FOR THIS CHAPTER.

3.1 F#

F# is a relatively young .NET-based language, first released in 2003. It is a strongly-typed multiple-paradigm language, with a syntax that is primarily functional, resembling OCaml. Although F# is not as widely used as C#, Java and the like, it is currently experiencing increasing adaptation among developers[?]. Besides supporting multiple paradigms and a reasonable subset of functional language features (such as pattern matching), F#'s primary strength is its interoperability with the rest of the .NET ecosystem. Like C#, F# programs are compiled into Microsoft's Common Intermediate Language, and executed using Microsoft's Common Language Runtime.

Therefore, F# programs have full access to the standard .NET library, just as it can also readily import and use classes and methods from arbitrary C# libraries.

For FShark, F# has been selected as a source language for several reasons. First, most of F#'s syntax is readily translatable into Futhark syntax, as long as the programmer stays away from using any of F#'s non-functional constructs, like `async` or its object oriented features. Second, as F# effortlessly interoperates with C# programs, and C# has plenty of OpenCL libraries available, we can write imperative OpenCL-powered programs in C#, for use in F# projects.

3.2 Futhark

Quoting from Futhark's own homepage,

Futhark is a small programming language designed to be compiled to efficient parallel code. It is a statically typed, data-parallel, and purely functional array language in the ML family, and comes with a heavily optimising ahead-of-time compiler that presently generates GPU code via OpenCL, although the language itself is hardware-agnostic.

So far, plenty of handwritten GPU benchmark programs implemented in CUDA et al, has been ported to Futhark, with significant performance gains as a result. [?]. With these results in mind, it makes sense to start implementing other parallelizable algorithms and programs in Futhark. However, in the grand scheme of things, Futhark is still a relatively obscure programming language, and is almost solely used in academic settings.

With Futhark being a purely functional programming language, it has very few imperative language constructs available, and the few that it has, like in-place updates, are merely syntactic sugar for other existing library function calls.

As Futhark's main functionality is generating OpenCL kernels, it is in principle possible to compile Futhark programs for any language that are able to interface with the OpenCL API.

As a target language for F# translations, Futhark is ideal as we can identify and relatively easily translate a subset of the F# language to equivalent Futhark code, as the syntax itself is very similar. Even though F# also allows plenty of imperative and object oriented programming, FShark blocks the user from using these constructs, by failing at FShark compile time.

3.3 C#

C# C# C# C#

Chapter 4

The **FShark** language

t	$::=$	$\text{int8} \mid \text{int16} \mid \text{int} \mid \text{int64}$	(Integers)
	$ $	$\text{uint8} \mid \text{uint16} \mid \text{uint} \mid \text{uint64}$	(Unsigned integers)
	$ $	$\text{single} \mid \text{double}$	(Floats)
	$ $	bool	(Booleans)
	$ $	$(t_0 * \dots * t_n)$	(Tuples)
	$ $	$\{\text{id}_0 : t_0; \dots; \text{id}_n : t_n\}$	(Records)
	$ $	$t_0 \text{ array}$	(Arrays)
k	$::=$	n	(Integer)
	$ $	f	(Float)
	$ $	b	(Boolean)
	$ $	(k_0, \dots, k_n)	(Tuple)
	$ $	$\{\text{id}_0 = k_0; \dots; \text{id}_n = k_n\}$	(Record)
	$ $	$[k_0; \dots; k_n]$	(Array)
p	$::=$	id	(Name pattern)
	$ $	(p_0, \dots, p_n)	(Tuple pattern)

Figure 4.1: The FShark syntax

e	$::=$	(e_0) k v (k_0, \dots, k_n) $\{\text{id}_0 = k_0; \dots; \text{id}_n = k_n\}$ $[k_0; \dots; k_n]$ $e_1 \odot e_2$ $-e$ $\text{not } e$ $\text{if } e_1 \text{ then } e_2 \text{ else } e_3$ $v.[e_0] \dots [e_n]$ $v.\text{id}$ $v_0.v_1$ $\text{let } p = e_1 \text{ in } e_2$ $v \ e_0 \dots e_n$ $v \ e_0 \dots e_n$
WHAT ABOUT LAMBDAS?		
fun	$::=$	$[\text{<FSharkEntry>}] \text{ let } v (v_1 : t_1) \dots (v_n : t_n) : t = e$ $\text{let } v (v_1 : t_1) \dots (v_n : t_n) : t' = e,$ $(\text{for any } i \in 1..n, t_i \text{ is not a tuple})$
typealias	$::=$	$\text{type } v = t$
module	$::=$	$\text{module } v = \text{prog}' \ \text{progs}'$
prog	$::=$	$\text{module } \text{prog}$ $\text{prog}' \ \text{prog}$ ϵ
prog'	$::=$	typealias
		fun
progs'	$::=$	$\text{prog}' \ \text{progs}'$
		ϵ

Figure 4.2: The FShark syntax, expressions

4.0.1 F# operators available in FShark

The F# subset chosen for FShark is described in this subsection. Note that all

Arithmetic operators

The set of supported arithmetic operators is addition (+), binary subtraction and unary negation (-), multiplication (*), division (/) and modulus (%).

Boolean operators

FShark currently supports logical AND (&&), logical OR (||), less- and greater-than (<, >), less- and greater-or-equal (<=, >=), equality (=), inequality (<>) and logical negation (not).

Special operators

FShark also supports some of F#'s syntactic sugar. These operators might not have direct Futhark counterparts, but their applications can be rewritten in Futhark for equivalent functionality. The supported operators are back- and forward pipes (<| and |>), and the range operator ($e_0 \dots e_1$), which generates the sequence of numbers in the interval $[e_0, e_1]$. Note that in FShark, the range operator must be used inside an array as so `[|e0..1|]` so we adhere to using arrays and not lists in our FShark programs.

Figure 4.3: FShark operators

of these operators are overloaded and defined for all integer and floating point types in F#.

4.0.2 F# standard library functions available in FShark

FShark supports a subset of the F# standard library. These are functions that are imported in F# modules by default.

Currently, bitwise operators like bitwise-AND and bitwise-OR are missing, but they should be relatively simple to add to the FShark subset, by adding them to the set of supported operators in the FShark compiler.

4.0.3 On the F# subset selected for FShark

For selecting the F# subset to support in FShark, I chose to look at what functions that were included in F#'s prelude. That is, the functions that are available in an F# program without having to open their containing module first. Fortunately, F# opens several modules by default of which I only needed to look in two different ones, to be able to support a reasonable amount of F# built-ins in FShark.

The primary module used in my supported F# subset is the module `FSharp.Core.Operators`. This module contained not only the standard arithmetic described in figure 4.3, but also most¹ of the functions shown in the figure 4.4. Except for unit type

¹except for some conversion functions, found in `FSharp.Core.ExtraTopLevelOperators`

id

The identity function.

Common math function

The square root function (`sqrt`), the absolute value (`abs`), the natural exponential function (`exp`), the natural- and the decimal logarithm (`log` and `log10`).

Common trigonometric functions

Sine, cosine and tangent functions (both standard and hyperbolic): `sin`, `cos`, `tan`), `sinh`, `cosh` and `tanh`. Also one- and two-argument arctangent: `atan` and `atan2`.

Rounding functions

FShark supports all of F#s rounding functions: `floor`, `ceil`, `round` and `truncate`.

Number conversion functions

FShark supports all of F#s number conversion functions. For all the following functions t , $te = e'$, $e : t_0$, $e' : t$, barring exceptions like trying to convert a too large 64-bit integer into a 32-bit integer.

The conversion functions available are `int8`, `int16`, `int`, `int64`, `uint8`, `uint16`, `uint`, `uin64`, `single`, `double`, `bool`.

Various common number functions

`min`, `max`, `sign` and `compare`.

Figure 4.4: FShark operators

functions like `failwith`, `exit` and `async`, most of the functions and operators `FSharp.Core.Operators` have direct counterparts in Futhark's prelude, with equivalent functionality: All except for four of operators and functions chosen for FShark are in fact implemented in Futhark's `math.fut` library. It was therefore an obvious decision to support these functions and operators in FShark.

However, for the remaining four functions, that didn't have equivalents in Futhark's `math.fut`, their function calls are replaced with their identities instead. In example, whereas the FShark code

```
exp x
```

is written in Futhark as

```
exp x
```

because the `exp` function is also available in `math.fut`, the FShark code

```
cosh x
```

is rewritten as the full hyperbolic sine function instead, as so

```
((exp x) + (exp (-x))) / 2.0
```

$e ::=$	k	Constant
	v	Variable
	(k_0, \dots, k_n)	(Tuple expression)
	$\{\text{id}_0 = k_0; \dots; \text{id}_n = k_n\}$	(Record expression)
	$[k_0; \dots; k_n]$	(Array expression)
	$e_1 \odot e_2$	(Binary operator)
	$-e$	(Prefix minus)
	$\text{not } e$	(Logical negation)

Figure 4.5: f binary operators

$e ::=$	k	Constant
	v	Variable
	(k_0, \dots, k_n)	(Tuple expression)
	$\{\text{id}_0 = k_0; \dots; \text{id}_n = k_n\}$	(Record expression)
	$[k_0; \dots; k_n]$	(Array expression)
	$e_1 \odot e_2$	(Binary operator)
	$-e$	(Prefix minus)
	$\text{not } e$	(Logical negation)

Figure 4.6: FShark SOACs

These rewritings are not pretty to look at from a programmer’s perspective, but FSharks Futhark code is not meant to be read by humans anyhow.

(MAYBE INVESTIGATE WHETHER INLINING THESE HAS PERFORMANCE PENALTIES)

4.0.4 The correctness of the FShark subset.

When transpiling code from one language to another, it is absolutely vital that the programmer can trust, that the resulting code in the target language is semantically equivalent to the source code. In FSharks case, it means that any program written using the FShark subset, must have the same result no matter whether it is run natively as F# code, or it is run as FShark compiled Futhark code.

I.e., one could imagine a programming language which had defined the function `log` not as the natural logarithm, but instead the binary logarithm. In such a case, the translation from that language to Futhark would still go without a hitch, and without any type errors to hint at the impending catastrophe. However, the native result with the Futhark result would be wildly different.

To ensure that every operator and function in the FShark subset has equivalent results, no matter whether the FShark code is run as native F# code, or compiled into Futhark, I have written a test suite with unit tests for each element in the F# subset.

This is described further in chapter ??

(THESE ARE NOT ACTUALLY DONE YET) (Are unit tests enough?)

LANGUAGE REFERENCE

4.1 FSharkPrelude

Besides defining an F# subset suitable for Futhark translation, it was also imperative to create a library of SOACs and array functions for FShark, to make it possible to write programs with parallel higher-order array functions.

Similarly to how the subset of math functions chosen from F# to include in the FShark was chosen, the SOACs and array function included in the FSharkPrelude has been picked directly from the Futhark libraries `futlib/array.fut` and `futlib/soacs.fut`. The FSharkPrelude doesn't discriminate between array functions and SOACs, as maintaining and importing two different prelude files in FShark was needlessly complicated.

The FSharkPrelude consists of functions which are directly named after their Futhark counterparts, and have equivalent functionality. This prelude, together with the FShark subset, is what makes up the FShark language. When FShark developers are writing modules in FShark, they are guaranteed that their FShark programs has the same results, no matter whether their programs are executed like native F# code, or compiled and executed as Futhark.

The FSharkPrelude versions of Futhark functions are defined in three different ways.

Functions like the SOAC `map` and the array function `length` have direct F# equivalents, and are therefore implemented as calls to `Array.map` and `Array.length` respectively. For `map` for example, we have the following definition:

```
module FSharkPrelude =  
  ...  
  
  let Map f aa =  
    Array.map f aa  
  
  ...
```

Some Futhark SOACs, like `reduce`, takes a neutral element as one of the arguments in their function calls, whilst their F# counterparts (`Array.reduce`) does only take an operator and an array as arguments. To define the FShark SOAC so that it is equivalent to the Futhark version, it has been defined as so:

```
module FSharkPrelude =  
  ...  
  
  let Reduce (op: 'a -> 'a -> 'a) (neutral : 'a) (xs : 'a array) =  
    let xs' = Array.append [|neutral|] xs  
    in Array.reduce op xs'  
  
  ...
```

Other functions, like the `map` functions which takes multiple arrays as arguments, require a bit of assembly first. For those `map` functions, we zip the arguments before using `Array.map` as usual:

```

module FSharkPrelude =
  ...
  let Map5 f aa bb cc dd ee =
    let curry f (a,b,c,d,e) = f a b c d e
    let xs = Zip5 aa bb cc dd ee
    in Array.map (curry f) xs
  ...

```

Lastly, some functions does not have F# counterparts. In example, we implement `scatter` using a for-loop:

```

module FSharkPrelude =
  ...
  let Scatter (dest : 'a array) (is : int array) (vs : 'a array) :
    for (i,v) in Zip is vs do
      dest.[i] <- v
    dest
  ...

```

The complete list of available SOACs and array functions is available in appendix ??.

Note that calls to `FSharkPrelude` functions are caught and exchanged for Futhark functions during the `FShark` compilation, as described in sec ??.

Several of Futhark's SOACs, such as `map`, already has F# versions that are directly equivalent.

But there are several issues with just letting the `FShark` programmer use But many of these F# functions are contained in `HER KOMMER DER MERE`

Chapter 5

The **FShark** Compiler and Wrapper

Introduction

Parsing and building a regular F# program is trivial when using official build tools like `msbuild` or `fsharpc`. But in the case of **FShark**, we are not interested in the final result from the F# compiler, but merely its half-finished product.

As the F# Software Foundation offers the official F# Compiler as a freely available NuGet package for F# projects, we can use this package `FSharp.Compiler.Services` to parse the entire input **FShark** program and give us a Typed Abstract Syntax Tree of the **FSharp** expressions therein.

FIGURE HERE OF THE FSHARP COMPILATION

Figure 5.1: The **FShark** compilation pipeline

The F# Software Foundation actively encourages developers to create projects using the F# compiler library, they have published the collected F# compiler as a NuGet package, alongside a tutorial??on the usage of the various compiler parts.

For **FShark**, the Compiler Services package is used to compile a Typed Abstract Syntax Tree from a wellformed **FShark** source code file, which we then convert into- and print as a valid **Futhark** program. The Typed Abstract Syntax Tree is merely an AST that already has tagged all the contained expressions with their respective types.

We'll start with a detailed explanation of the **FShark** Compiler Pipeline.

5.0.1 The **FShark** Compiler Pipeline in practice

To examine the compiler pipeline in action, we'll go through the motions with the small example program displayed in figure 5.2.

We begin by constructing an instance of the **FSharkWrapper**. It has the following mandatory arguments:

```

1  module FSharkExample
2  open FShark.Main
3
4  [<EntryPoint>]
5  let main argv =
6      let wrapper =
7          new FSharkWrapper (
8              libName="ExampleModule",
9              tmpRoot="/home/mikkel/FShark",
10             preludePath= "/home/mikkel/Documents/fshark/F-
               ↪ SharkPrelude/bin/Debug/FSharkPrelude.dll",
11             openCL=true,
12             unsafe=true,
13             debug=false
14         )
15     wrapper.AddSourceFile "../srcs/ExampleModule.fs"
16     wrapper.CompileAndLoad
17     let xs = [|1;2;3;4|]
18     let input = [|xs|] : obj array
19     let xs' = wrapper.InvokeFunction "MapPlusTwo" input :?>
               ↪ int array
20     printfn "Mapping (+2) over %A gives us %A" xs xs'
21     0

```

Figure 5.2: An F# program using FShark

libName

This is the library name for the FShark program. In the final Futhark .cs and .dll files, the main class will have the same name as the libName. This doesn't really matter if FShark is just used as a JIT compiler, but it's good to have a proper name if the user only wants to use the compiler parts of FShark.

tmpRoot

The FShark compiler works in its own temporary directory. This argument must point to a directory where F# can write files and execute subprocesses (Futhark- and C# compilers) which also has to write files.

preludePath

The FShark compiler needs the FShark prelude available to compile FShark programs.

openCL

Although Futhark (and therefore FShark) is most effective on OpenCL-enabled computers, the benchmarks in ?? still show a significant speed increase for non-OpenCL Futhark over native F# code. Therefore, FShark is also available for non-OpenCL users. Use this flag to tell FShark whether Futhark should compile C# with or without OpenCL.

unsafe

For some Futhark programs, the Futhark compiler itself is unable to tell whether certain array operations or SOAC usages are safe, and will stop the compilation, even though the code should (and does) indeed work. To enable these unsafe operations, pass a `true` flag to the compiler.

debug

Passing the `debug` flag to the FShark compiler enables various runtime debugging features, for instance benchmarking the time it takes to run various parts of the compiler.

Now, we can pass a source file to the FShark wrapper, compile¹ it and load it into the FShark wrapper object.

To use the compiled FShark function, we must first wrap our designated input in an `obj array`. In this case, our chosen FShark function takes one argument, an `int array`. We define this array, and construct an argument array containing this single element. If the FShark function takes two arguments, we define an input `obj array` with two elements, and so forth. It is important to declare the input array as an `obj array`. Otherwise, F#'s own type checker might very well faultily infer the input array as something else. In this particular case, `input` would've been inferred as being an `int array array`, until we declared its type specifically.

We then invoke the desired function through the wrapper. As all reflection-invoked functions return a value of type `obj`, we need to downcast this object manually. In this example, we use F#'s downcast operator `(:?)` to declare the return value as an `int array`. The actual return type is always the same as the return type declared in the source FShark file.

5.0.2 When FShark Wrapper Compiles

The general way to compile and load an FShark program into the FShark Wrapper, is by adding FShark source files to the wrapper object by calling the `AddSourceFile` method, and followingly calling the `CompileAndLoad` method. Although the FShark wrapper also offers other methods of loading and compilation, this is the primary one, as it initiates the entire FShark compilation pipeline.

When calling `CompileAndLoad`, the supplied FShark source files are concatenated into one long source file, and written to a temporary location. An `FSharpChecker` is then initialized, so we can parse and type check the concatenated source code. The `FSharpChecker` is a class exported by the FSharp Compiler Services, and is a class that lets developers use part of the F# compilation pipeline at runtime.

We supply the `FSharpChecker` with the path to our precompiled `FSharkPrelude` assembly, and then call its `ParseAndCheckProject` method on to receive an assembly value, which contains the complete Typed Abstract Syntax Tree of our FShark program, in the form of an `FSharpImplementationFileDeclaration`.

If the FShark developer followed the guidelines to write a well-formed FShark module, the main declaration of the program, the `FSharpImplemen-`

¹See subsection ??

tationFileDeclaration, should contain a single FSharpEntity, which in turn contains all the remaining declarations in the program.

The declaration types within F#'s Typed AST

The FSharpImplementationFileDeclaration type has three union cases.

InitAction of FSharpExpr

InitActions are FSharpExprs that are executed at the initialization of the containing entity. These are not supported in FShark.

Entity of FSharpEntity * FSharpImplementationFileDeclaration list

An Entity is the declaration of a type or a module. In the case of FShark, three different kinds of entities are supported:

FSharpRecords are standard record types, and can be translated to Futhark records with ease. This entity has an empty FSharpImplementationFileDeclaration list.

FSharpAbbreviations are type abbreviations, and are easily translated into Futhark type aliases. This entity has an empty FSharpImplementationFileDeclaration list.

FSharpModules are named modules which contains subdeclarations. In this case, we retrieve the subdeclarations from the FSharpImplementationFileDeclaration list. The FShark compiler supports building FShark modules, but current limitations demands that modules are flattened when compiled to Futhark. This also means that function name prefixes in function calls are stripped when compiled to Futhark.

MemberOrFunctionOrValue of FSharpMemberOrFunctionOrValue * FSharpMemberOrFunctionOrValue

F# doesn't differ between functions and values, which means that a function is merely a value with arguments. A pattern matched MemberOrFunctionOrValue value has the form MemberOrFunctionOrValue (v, args, exp), where v contains the name and the type of the variable. If the args list is empty, v is simply a variable. If not, v is a function. exp contains the FSharpExpr that v is bound to. An FSharpExpr can be anything from a numeric constant to a very long function body.

In figure 5.3 we see a small but valid FShark program. It reads like a regular F# program, but contains the three vital parts that makes it usable as an FShark program.

- The module declaration on the first line declares that the following code is inside a module. In this case, we are declaring the module ExampleModule, although we could use any valid F# module name. As shown in figure 5.4, the top module declaration falls away during compilation, so only the top module contents are left.
- This open statement ensures that the F# Compiler Services has access to the FSharkPrelude during the compilation. It is possible to write an

```

1  module ExampleModule
2  open FSharkPrelude
3
4  module SomeValues =
5      let Four : int = 4
6
7      let SomePlus (x : int) (y : int) : int = x + y
8
9      [<FSharkEntry>]
10     let TimesTwo (x : int) : int =
11         SomeValues.SomePlus x x
12
13     [<FSharkEntry>]
14     let MapPlusTwo (xs : int array) : int array =
15         Map ((+)2) xs
16
17     let PlusSeven (x : int) : int =
18         SomeValues.SomePlus x 7

```

Figure 5.3: A valid FShark program

FShark program which doesn't use the FSharkPrelude, but this removes access to the SOACs that we use to write our data parallel programs.

- The [<FSharkEntry>] attributed function TimesTwo ensures that the resulting Futhark library from the FShark compiler has at least one entry point function. Without any entry point functions, we won't have any functions in the final compiled FShark program.

In figure 5.4 we see the resulting Futhark program. For now, we will ignore the transformations that have happened, except for two things: The Map function (called from FSharkPrelude) has been rewritten as the plain Futhark SOAC map in lowercase, and the module SomeValues has been flattened (see sec ?? for future plans.)

```

let Four : i32 = 4i32
let SomePlus (x : i32) (y : i32) : i32 =
    ((x i32.+ y))
entry TimesTwo (x : i32) : i32 =
    unsafe SomePlus(x) (x)
entry MapPlusTwo (xs : []i32) : []i32 =
    unsafe map (let x = 2i32 in
        (\(y : i32) -> ((x i32.+ y)))) (xs)
let PlusSeven (x : i32) : i32 =
    SomePlus(x) (7i32)

```

Figure 5.4: A valid FShark program, compiled to Futhark

$\llbracket Int8 \rrbracket$	=	Prim Int FInt8	$\llbracket Int16 \rrbracket$	=	Prim Int FInt16
$\llbracket Int32 \rrbracket$	=	Prim Int FInt32	$\llbracket Int64 \rrbracket$	=	Prim Int FInt64
$\llbracket UInt8 \rrbracket$	=	Prim UInt FUInt8	$\llbracket UInt16 \rrbracket$	=	Prim UInt FUInt16
$\llbracket UInt32 \rrbracket$	=	Prim UInt FUInt32	$\llbracket UInt64 \rrbracket$	=	Prim UInt FUInt64
$\llbracket Single \rrbracket$	=	Prim Float FSingle	$\llbracket Double \rrbracket$	=	Prim Float FDouble
$\llbracket Boolean \rrbracket$	=	Prim Bool	$\llbracket \tau \rrbracket$	=	FSharkArray $\llbracket \tau \rrbracket$
$\llbracket (\tau_0 \times \dots \times \tau_n) \rrbracket$	=	FSharkTuple $\llbracket \tau_0 \rrbracket, \dots, \llbracket \tau_n \rrbracket$			

INSERT NOTE ON RULE FOR TUPLE ('a [] * long [])

Figure 5.5: F# (.NET) types to FSharkIL types

$\llbracket PrimIntFInt8 \rrbracket$	=	i8	$\llbracket PrimIntFInt16 \rrbracket$	=	i16
$\llbracket PrimIntFInt32 \rrbracket$	=	i32	$\llbracket PrimIntFInt64 \rrbracket$	=	64
$\llbracket PrimUIntFUInt8 \rrbracket$	=	u8	$\llbracket PrimUIntFUInt16 \rrbracket$	=	u16
$\llbracket PrimUIntFUInt32 \rrbracket$	=	u32	$\llbracket PrimUIntFUInt64 \rrbracket$	=	u64
$\llbracket PrimFloatFSingle \rrbracket$	=	f32	$\llbracket PrimFloatFDouble \rrbracket$	=	f64
$\llbracket PrimBool \rrbracket$	=	bool	$\llbracket \tau \rrbracket$	=	FSharkAr
$\llbracket (\tau_0 \times \dots \times \tau_n) \rrbracket$	=	FSharkTuple $\llbracket \tau_0 \rrbracket, \dots, \llbracket \tau_n \rrbracket$			

Figure 5.6: FSharkIL types to Futhark types

This Futhark program is then stored in a temporary location in the user's file system, and compiled into as a library, using Futhark's C# compiler, either with or without OpenCL support. Finally after this compilation, we can invoke the resulting .dll file from within the FShark-using F# program.

5.0.3 Building **FShark** from the Typed AST

FShark supports just a subset of the F# language, which also means that only a subset of F#'s FSharpExpr

Only the supported FSharpExpr's has been listed here, but the full range of FSharpExpr's are available on [?].

5.0.4 FSharp-to-FSharkIL rules

INTRODUCTION HERE

For these translations, we will disregard that all FSharpExprs are union cases of the F# data type BasicPatterns.

$\llbracket \text{Const}(\text{obj}, \tau) \rrbracket$	$= \text{Const}(\text{obj}, \llbracket \tau \rrbracket)$
$\llbracket \text{Value}(v) \rrbracket$	$= \text{Var}(v)$
$\llbracket \text{AddressOf}(v) \rrbracket$	$= \llbracket v \rrbracket$
$\llbracket \text{NewTuple}(_, [e_0, \dots, e_1]) \rrbracket$	$= \text{Tuple}(\llbracket [e_0], \dots, [e_1] \rrbracket)$
$\llbracket \text{NewRecord}((v_0 : \tau_0 * \dots * v_n : \tau_n), [e_0, \dots, e_1]) \rrbracket$	$= \text{Record}(\llbracket (v_0, [e_0]), \dots, (v_n, [e_n]) \rrbracket)$
$\llbracket \text{NewArray}(\tau, [e_0, \dots, e_1]) \rrbracket$	$= \text{List}(\llbracket \tau \rrbracket, \llbracket [e_0], \dots, [e_1] \rrbracket)$
$\llbracket \text{TupleGet}(_, i, e) \rrbracket$	$= \text{TupleGet}(\llbracket e \rrbracket, i)$
$\llbracket \text{FSharpFieldGet}(e, _, \text{field}) \rrbracket$	$= \text{RecordGet}(\text{field}, \llbracket e \rrbracket)$
$\llbracket \text{Call}(_, \text{GetArray}, _, \text{nil}, [e_0, e_1]) \rrbracket$	$= \text{ArrayIndex}(\llbracket e_0 \rrbracket, \llbracket [e_1] \rrbracket)$
$\llbracket \text{Call}(_, \text{name}, _, \text{nil}, [e_0, \dots, e_n]) \rrbracket$	$= \text{Call}(\text{name}, \llbracket [e_0], \dots, [e_n] \rrbracket)$
$\llbracket \text{Call}(_, \text{name}, _, \tau, [e_0, \dots, e_n]) \rrbracket$	$= \text{TypedCall}(\llbracket \tau \rrbracket, \text{name}, \llbracket [e_0], \dots, [e_n] \rrbracket)$
$\llbracket \text{Call}(_, \text{infixOp}, _, \tau, [e_0, e_1]) \rrbracket$	$= \text{InfixOp}(\text{infixOp}, \llbracket \tau \rrbracket, \llbracket [e_0], [e_1] \rrbracket)$
$\llbracket \text{Call}(_, \text{unaryOp}, _, \tau, [e_0]) \rrbracket$	$= \text{UnaryOp}(\text{unaryOp}, \llbracket \tau \rrbracket, \llbracket [e_0] \rrbracket)$
$\llbracket \text{Let}(v, e_0, e_1) \rrbracket$	$= \text{LetIn}(v, \llbracket [e_0], [e_1] \rrbracket)$
$\llbracket \text{IfThenElse}(e_0, e_1, e_2) \rrbracket$	$= \text{If}(\llbracket [e_0], [e_1], [e_2] \rrbracket)$
$\llbracket \text{Lambda}((v : \tau), e) \rrbracket$	$= \text{Lambda}(v, \llbracket \tau \rrbracket, \llbracket [e] \rrbracket)$
$\llbracket \text{Application}(\text{func}, _, [e_0, \dots, e_n]) \rrbracket$	$= \text{Application}(\llbracket [func] \rrbracket, \llbracket [e_0], \dots, [e_n] \rrbracket)$
$\llbracket \text{TypeLambda}(e) \rrbracket$	$= \llbracket [e] \rrbracket$
$\llbracket \text{DecisionTree}(_, _) \rrbracket$	$= \text{Pass}$
$\llbracket \text{DecisionTreeSuccess}(_, _) \rrbracket$	$= \text{Pass}$

Figure 5.7: FSharpExpr expressions to FSharkIL

$\llbracket \text{Const}(\text{obj}, \tau) \rrbracket$	$= \text{obj} \llbracket \tau \rrbracket$
$\llbracket \text{Var}(v) \rrbracket$	$= v$
$\llbracket \text{Tuple}([e_0, \dots, e_n]) \rrbracket$	$= ([e_0], \dots, [e_n])$
$\llbracket \text{Record}([(v_0, e_0), \dots, (v_n, e_n)]) \rrbracket$	$= \{v_0 = [e_0], \dots, v_n = [e_n]\}$
$\llbracket \text{List}(\llbracket \tau \rrbracket, \llbracket [e_0], \dots, [e_n] \rrbracket) \rrbracket$	$= ([e_0], \dots, [e_n])$
$\llbracket \text{TupleGet}(\llbracket [e] \rrbracket, i) \rrbracket$	$= [e].i$
$\llbracket \text{RecordGet}(\text{field}, e) \rrbracket$	$= [e].\text{field}$
$\llbracket \text{ArrayIndex}(e_{\text{arr}}, [e_0, \dots, e_n]) \rrbracket$	$= [e_{\text{arr}}] \llbracket [e_0], \dots, [e_n] \rrbracket$
$\llbracket \text{Call}(\text{name}, [e_0, \dots, e_n]) \rrbracket$	$= \text{name}([e_0]) \dots ([e_n])$
$\llbracket \text{TypedCall}(\llbracket \tau \rrbracket, \text{name}, \llbracket [e_0], \dots, [e_n] \rrbracket) \rrbracket$	$= \llbracket \tau \rrbracket.\text{name}([e_0]) \dots ([e_n])$
$\llbracket \text{InfixOp}(\text{infixOp}, \llbracket \tau \rrbracket, \llbracket [e_0], [e_1] \rrbracket) \rrbracket$	$= ([e_0]) \llbracket \tau \rrbracket.\text{infixOp}([e_1])$
$\llbracket \text{UnaryOp}(\text{unaryOp}, \llbracket \tau \rrbracket, \llbracket [e_0] \rrbracket) \rrbracket$	$= \llbracket \tau \rrbracket.\text{unaryOp}([e_0])$
$\llbracket \text{LetIn}(v, [e_0], [e_1]) \rrbracket$	$= \text{let } v = [e_0] \text{ in } [e_1]$
$\llbracket \text{If}(\llbracket [e_0], [e_1], [e_2] \rrbracket) \rrbracket$	$= \text{if } [e_0] \text{ then } [e_1] \text{ else } [e_2]$
$\llbracket \text{Lambda}(v, \llbracket \tau \rrbracket, \llbracket [e] \rrbracket) \rrbracket$	$= \lambda(v : \llbracket \tau \rrbracket) \rightarrow [e]$
$\llbracket \text{Application}(\llbracket [func] \rrbracket, \llbracket [e_0], \dots, [e_n] \rrbracket) \rrbracket$	$= ([func])([e_0]) \dots ([e_n])$
$\llbracket \text{Pass} \rrbracket$	$= \epsilon$

Figure 5.8: FSharkIL expressions to Futhark

Chapter 6

The Futhark C# backend

Chapter 7

Evaluation

Chapter 8

Current limitations

Chapter 9

Method

Chapter 10

Related work

Chapter 11

Future work

Chapter 12

Conclusion

Chapter 13

Array handling in FShark

F# is a functional programming language on top of the .NET framework, which means that its primitive types like `int`, `float` and `list` all correspond to already existing classes in the .NET framework. In example, F#'s `int` is an alias for .NET's `System.Int32` and `float` is an alias for `System.Double`.

Therefore, we also find corresponding .NET classes for both F# lists and arrays. Lists are `FSharp.Collections.FSharpList`s, and arrays are `System.Array`. (Note that `FSharpList` is available from any .NET framework language, as long as the corresponding assembly is referenced).

Although it is common to use lists in functional programs, the F# subset covered by FShark does not include lists – In Futhark, and therefore also FShark, our main goal is not handling list elements one after another, but rather parallelizing computations across entire arrays of data simultaneously.

The `FSharp.Collections.FSharpList` is implemented as a singly-linked list. SOACs called on singly-linked lists are inherently unparallelizable, as the SOACs must traverse the list sequentially. In example, calling `map f` on a singly linked list `(x : xs)` means computing `f x` and inserting the result into `map f xs` recursively. We can do some parallel computations for these kinds of SOACs, i.e. by making the main thread traverse the list and spawn a thread for each element computation. However we will still have suboptimal memory access performance, as the elements in the singly linked list doesn't have any guarantees regarding their location in RAM, which means we are going to perform many more memory loads compared to if we were performing calculations on elements in a sequentially stored array elements in RAM.

Chapter 14

FSharks interoperability between F# and Futhark (C#)

FShark stands on three legs: The FShark compiler, the Futhark C# code generator, and the FSharkWrapper. The compiler is responsible for compiling FShark source code into Futhark source code, and the C# code generator takes the result Futhark source code, and compiles OpenCL powered C# libraries, which can be imported directly back into F#.

It is of course possible to use the compiler and the code generator as individual modules, but for this project, the FSharkWrapper has been designed to let users use FShark without having to understand any of the underlying pipeline.

To illustrate this; take a look at figure ???. In the first line, the user initializes the FSharkWrapper with the arguments necessary to use the wrapper itself. In the second line, the user adds a source file to the wrapper by its path. In the third line, the user tells the wrapper to run the compilation pipeline. Assuming that the compilation goes well, the user can then invoke some function from the FShark program in line four.

Here, calling the `CompileAndLoad()` function triggers the entire FShark pipeline as described in ??, and does then have a function available for the user to call afterwards.

However, as this is the default way of using FShark, we are currently calling `CompileAndLoad()` every time we use the FShark program. This is happening even though we only need the final compiled C# assembly to load back into F# at runtime.

Everytime we run the FShark compiler pipeline, we are therefore also

1. parsing, typechecking and generating a TAST from the FShark code, using FSharp's compiler.
2. generating Futhark source code from the FSharp TAST
3. Writing the Futhark source code to disk
4. running the Futhark compiler and C# code generator on the Futhark source code

5. running the mono C# compiler on the resulting C# source code

For two selected benchmarks we have the following times

14.0.1 Pros and cons of the current design

As there are demonstratively great performance gains to be won by only using the compiler part of the FShark pipeline, it is worth discussing whether the rest of the FShark pipeline should remain.

Besides eliminating the entire compilation operation at every FShark execution, a compiler-only approach to the FShark compiler would give us the following advantages:

- **Standalone-modules first:** As the compiler is now only used once, the resulting Futhark assembly is readily importable in any .NET project, as long as the required Mono libraries are also available. This goes not only for the user who just compiled the assembly, but also for any other user who has acquired the necessary Mono libraries. This means that the FShark developer can use and share the FShark assemblies with colleagues and coworkers like any other sharable .NET library, as this is indeed what a compiled Futhark C# library amounts to.

Corollarily; this FShark design would make FShark is useful for generating high-performing .NET libraries. (Although one could write such libraries in Futhark instead of FShark.)

- **Static typing of FShark module:** The current runtime-only approach means that the user must rely on reflection to call FShark functions. In this situation, all modern IDE comforts like autocompletion, and especially static type checking and inference falls away. For the following example , the current design demands that we first wrap our arguments in an `obj` array, before invoking the function `foo`. Furthermore, we must also downcast the result using F#'s downcasting operator `:?>`. Because we are upcasting our arguments to an `obj` array, we can actually pass any (correctly casted) array as an argument to our reflection-invoked function, without triggering any type errors at compile time. The same goes for the downcasted result from the function. We can cast the result as whatever type we like, and not run into any trouble until we finally run the compiled program. However, if we use FShark to generate assemblies instead, we now have all the type information available at compile time. Our compiler will block us from compiling the program by giving us useful type errors. Last but not least, we can remove all the casting operations that are littering the program.

- **Getting rid of, or at least trimming down, the FSharkWrapper:**

However, the current design of FShark also has some advantages that follows automatically from the design.

- **Rapid FShark code development:** Currently, it is recommended that any FShark code for a project is also built as part of the project. By including the FShark file in the original project's source list, we can call the FShark module natively, without running the FShark compiler, to

prototype and debug the FShark code directly in our IDE, before we switch to using the compiled version of the FShark code.

- En mere

14.1 The future of **FShark** interoperability

With these considerations in mind, my future work on FSharks interoperability consists of reducing FSharkWrapper in size, so it only takes an FShark source path and a .NET assembly outpath as inputs, and does nothing more than orchestrating the FShark-, the Futhark- and the C#compiler. The current design is too complex, largely from supporting too many superfluous features like concatenating multiple sources, and so on.

I will also be researching the optimal way to keeping the FShark module development as close to the rest of the FShark-using project as possible, without

The current design enables direct prototyping, which must of course be kept in later versions of FShark.